

DIPLOMARBEIT

Localisation via ML Methods

Ausgeführt im Schuljahr 2019/20 von:

Algorithm	5AHIF
Ida Höningmann	
System Engineering	5AHIF
Peter Kain	

Betreuer / Betreuerin:

MMag. Dr. Michael Stifter

Wiener Neustadt, am November 18, 2019/20

Abgabevermerk:

Übernommen von:

Eidestattliche Erklärung

Hiermit erkläre ich an Eides statt, dass ich die vorliegende Arbeit selbstständig und ohne fremde Hilfe verfasst und keine anderen als die im Literaturverzeichnis angegeben Quellen und Hilfsmittel verwendet habe. Insbesondere versichere ich, dass ich alle wörtlichen und sinngemäßen Übernahmen aus anderen Werken als solche kenntlich gemacht habe.

Wiener Neustadt am November 18, 2019/20

Verfasser / Verfasserinnen:

Ida HÖNIGMANN

Peter KAIN

Contents

Eidestattliche Erklärung	i
Acknowledgement	iv
Kurzfassung	v
Abstract	vi
1 Introduction	1
1.1 Goal	1
1.2 Motivation	2
1.3 Outlook / Perspective	2
1.4 Equipment	3
2 Study of Literature	5
2.1 Different Approaches to the Problem	5
2.2 Depth perception	5
2.2.1 Depth sensation	5
2.3 Stereo Camera	5
2.3.1 Distortion	5
2.3.2 Image Rectification	6
2.3.3 Disparity Map	8
2.3.4 3d point cloud	8
2.4 LIDAR	8
2.5 Structure from Motion	8
2.6 Feature Tracking	8
3 Methodology	9
3.1 Challenges in the use of Stereo Cameras	9
3.2 Generating data	9
3.2.1 Blender	9
3.3 Image preprocessing	11
3.4 Neural Network	13
3.4.1 Convolutional Neural Network	13

3.4.2	Loss Function	13
3.4.3	Activation Function	15
3.5	Neural Network Implementations	17
3.6	TensorFlow	20
3.6.1	Computation graph	20
3.6.2	Alternatives to Tensorflow	21
4	ROS2	23
4.1	What is ROS?	23
4.1.1	Nodes	23
4.1.2	23
4.2	Why use ROS?	23
4.3	Comparing ROS and ROS2	23
5	Implementation	24
5.1	Generating test data	24
5.2	OpenCV	24
5.2.1	Greyscale	25
5.2.2	Resolution	27
5.2.3	Cropping	28
5.2.4	Saturated	29
5.2.5	Brightness	30
5.3	Neural Network	30
5.3.1	Structure of our Neural Network	30
5.4	C++ Implementation	31
5.5	Technical difficulties	31
6	Experiment 1	32
6.1	Environment	32
6.2	Setup	32
6.3	Sequence of Events	32
6.4	Results	32
7	Lessons learned	33
8	Experiment 2	34
8.1	Environment	34
8.2	Setup	34
8.3	Materials	34
8.4	Sequence of Events	34
8.5	Results	34
9	Conclusion	35

Acknowledgement

The authors would like to thank ...

Kurzfassung

asdf

Abstract

asdf

Chapter 1

Introduction

Author: Ida Höningmann

Robots are getting more and more mobile. While a few years ago their usage was mostly limited to aid factory automation, robots have found widespread adoption in a multitude of industries, such as self driving cars and autonomous delivery drones. A challenge frequently encountered is navigating in unknown environments, which either requires the robot to sense specific characteristics of its surroundings or to communicate with some external system.

The problem of navigation has been looked at from many different angles. One popular approach in mobile robotics is to use the GPS, an external positioning system. In order to determine the position of a robot using the GPS, it has to establish communication with at least four satellites. The exact position of each satellite as well as the current time is broadcast by the satellites. By measuring the time needed for the signal to reach the robot, the position can be calculated up to three meters accurately.

However, in some cases positioning a robot using external positioning methods is not possible. In the case of the GPS this can be due to obstacles interfering with the radio signals send by the satellites, for example occurring inside a building. In comparison this work focuses on a system that can navigate in outdoor as well as in indoor environments.

1.1 Goal

The goal of this diploma thesis is to implement a system which can localize a robot using no other sensors than a camera. This limitation was purposely chosen as the system will be used by future robotic students at the HTBLuVA (Technical Secondary College) and many robot systems used in the field of education are only poorly equipped with sensors that are able to detect its environment. One sensor used in the field of educational robotics is the either already equipped, or easily mountable camera.

As part of this thesis the authors not only want to implement an easy to use API for future robotic students, but to also show the possibilities and advantages of machine learning in localisation.

In order to accomplish precise localisation in various different surroundings, the authors plan on implementing a neural network. The neural network should take images, taken by the

camera, as an input, and outputs the relative distance to any object shown in the images. By using machine learning the system should be less dependent on a specific situation or setup in comparison to different camera based localisation methods. For example the localisation should work on objects varying in size and shape, as well as in different situations of lighting.

1.2 Motivation

In July 2019 the two authors of this work participated in the aerial tournament at the Global Conference on Educational Robotics held in Norman, Oklahoma. One of the two main challenges encountered at this tournament was landing a drone next to some randomly placed object, which colour, shape and size was known in advance.

The second challenge the participants at the tournament faced was flying from one side of randomly placed cardboard boxes to the other. The cardboard boxes, representing a mountain, are placed in one of various configurations, one of which can be seen in figure 1.1.




Figure 1.1: Seven cardboard boxes, representing a mountain, are placed in a random configuration. The team scores points if the drone passes the mountain while staying under the height limit indicated by the red dotted line.

The drones used at this aerial tournament are equipped with a camera, while lacking any other sensor that can be used to detect the obstacles and game items. Therefore the participants needed to be able to detect the distance to the object and the cardboard boxes using only the camera. At the Global Conference on Educational Robotics the authors of this work decided to detect the object based on its colour, but had to invest quite some time tweaking the values to get the localisation working correctly. Therefore the authors want to research and implement a method that is more robust than the colour based one.

1.3 Outlook / Perspective

The objective of this work is to create a system which uses machine learning methods in localising objects. After having trained the system, it should reliably return the x, y and z distances to an object, shown in two pictures taken from different angles. If this task turns

out to be too complex the system should be simplified by only returning one output number corresponding to distance from the drone to the object.

It is planned that the distance will be measured from the second camera position to the centre of object, as seen in Figure 1.2.




Figure 1.2: A drone tries to gather the data needed for the localisation of an object. After taking the first picture (grey drone with dotted lines) the drone flies to a second position (black drone with solid lines) to take another picture from a different point of view. The red line indicates the distance vector to be returned.

1.4 Equipment

Testing of the functionality of the localisation is to be performed on a Parrot Bebop 2. This drone was chosen as it is available to the authors during their time of research on this subject. This drone offers 1080p full HD video, which is believed to be more than sufficient for this project.

The Parrot Bebop 2 weighs 500g. The drone was not designed to carry any additional weight, therefore mounting a stereo camera, a LIDAR sensor or any other method of localisation using additional hardware is not recommended.

A wide-angle 14 megapixel fish-eye lens is mounted as a camera on the Parrot Bebop 2¹. It films 180 degrees vertically and horizontally and returns a 16:9 section which can be selected by specifying the vertical and horizontal angle of the virtual 16:9 camera. While the video resolution is limited to 1920 x 1080 pixels at 30 fps, the photo resolution is 4096 x 3072 pixels. A picture of the Parrot Bebop 2 can be seen in Figure 1.3.

¹Parrot Drones SAS. *Parrot Bebop 2 Specifications*. 2019. URL: <https://www.parrot.com/us/drones/parrot-bebop-2>.



Figure 1.3: The drone chosen for testing of the localisation system is a Parrot Bebop 2. It consists of four propellers each attached to a motor, a camera in the front, a battery located on the rear, some processors and a plastic frame to hold everything together.

It is assumed that the distance estimation using the system described in this thesis will work on various robotic systems. However, it will only be tested on this specific drone as other drones with a similar camera setup will behave similarly.

Chapter 2

Study of Literature

Author:

2.1 Different Approaches to the Problem

2.2 Depth perception

[TODO: humans, two eyes - gleich wie bei unserem Aufbau]

2.2.1 Depth sensation

[TODO: Pigeons, deer, children (visual cliff)]

2.3 Stereo Camera

The challenge of sensing distances to various objects has been solved using stereo vision cameras. Computer stereo vision systems use two horizontally displaced cameras to take two images which are then processed together to gather the information on the depth of the images. This process can be rather complicated as the distortions (more specifically the barrel distortion and the tangential distortion) of the images have to be undone, before the two images are projected onto a common plane, a disparity map can be created by comparison of the two images and a 3d point cloud can be generated from it. In most robotics applications this point cloud is then filtered in search of some object, which distance was sought-after.

2.3.1 Distortion

Barrel distortion occurs when the lens used by the camera has a higher magnification at the centre of the image than at the sides. This distortion can be visualized as seen in Figure 2.1. To undo this distortion the pixel values in an undistorted image have to be calculated based on the pixel values in the distorted image.

$$r_u = r_d(1 + kr_d^2) \quad (2.1)$$




Figure 2.1: The left shows the original image composed of straight horizontal and vertical lines. On the right image the effect of the barrel distortion can be perceived, which causes the lines to curve toward the outside of the image, causing the lines to appear in a barrel like shape.[TODO: change image to own]

describes the calculation which computes the distance from the centre in the undistorted image (r_u) based on the distance from the centre in the distorted image (r_d) and some distortion parameter k , which is specific to the lens used. Gribbon et.al. note in their work¹ that this rarely is an integer value, therefore different equations are proposed:

$$x_d = x_u M(k, r_u^2) \quad y_d = y_u M(k, r_u^2) \quad (2.2)$$

where the magnification factor $M(k, r_u^2)$ is

$$M(k, r_u^2) = \frac{1}{1 + k * M(k, r_u^2)^2 * r_u^2} \quad (2.3)$$

Tangential distortion in comparison to barrel distortion displaces points along the tangent of a circle placed at the centre of the image as seen in Figure 2.2.

The radius of the circle in Figure 2.2 is dependent on the point P . It can be calculated as the length between P and C . The length of the vector PP' is not uniform for all points and therefore depends on point P .

2.3.2 Image Rectification

Image Rectification projects multiple images taken from different points of view onto a common plane.

¹Donald G. Bailey K.T. Gribbon C.T. Johnston. “A real-time FPGA implementation of a barrel distortion correction algorithm with bilinear interpolation”. In: *Image and Vision Computing New Zealand*. 2003, pp. 408–413.




Figure 2.2: A point P is distorted along the tangent t of a circle placed at the middle of the image C with a radius r to a point P' . Distortions of this form are called tangential distortions.




Figure 2.3: Two images containing some point P are taken from the two points O_L and O_R . Point P is projected in the image planes as points P_L and P_R . E_L and E_R depict the epipoles.

Chan et. al. propose an image rectification algorithm², which follows this sequence of events:

1. At least seven matching points visible on both images are found.
2. The fundamental matrix (as well as the epipoles) are estimated.
3. The common region is identified (using epipolar geometry constraints).
4. The epipolar line is transferred and the Bresenham algorithm³ is used to extract pixel values.

²Hung Tat Tsui Zehzhi Chen Chengke Wu. “A new image rectification algorithm”. In: *Pattern Recognition Letters* 24.1-3 (2003), pp. 251–260.

³Jack Bresenham. “A linear algorithm for incremental digital display of circular arcs”. In: *Communications*

5. The rectified image is resampled.

2.3.3 Disparity Map

2.3.4 3d point cloud

[Input: Paper (gleiches Problem ohne NN) finden - Peter]

[Input: Video (eine Kamera, Entfernung zu Punkt (größe bekannt)) finden (ohne NN) - Peter]

[Vielleicht ist irgendetwas davon spannend: Links im Tex file]

2.4 LIDAR

2.5 Structure from Motion

2.6 Feature Tracking

of the ACM 20.2 (1977), pp. 100–106.

Chapter 3

Methodology

Author:

3.1 Challenges in the use of Stereo Cameras

Since many drones used in educational robotics can only carry a limited amount of weight it is not possible to attach a stereo camera to such a robot. Instead the functionality of a stereo camera on such a drone system can be mimicked by taking the first image, flying to a second position, located horizontally next to the first one and taking the second image. This process is not as precise as a stereo camera, where the two lenses are always positioned at an exact interval from one another. Therefore the output of this system might not work as reliable. Additionally other factors, such as differences in the two images due to some time passing between the taking of the images can have an effect on the accuracy.

Therefore the authors try to approach this challenge from the machine learning point of view. Neural networks can be taught to work with different changes in the environment and still return results with a superior quality to conventional implementation.

3.2 Generating data

Neural Networks require huge amounts of data to work reliably. Because of the author's limited time frame this test data will be generated with the help of Blender 2.8. Blender is a free program for designing and animating 3D objects, which also supports scripting with python to add or remove objects from a scene. The authors will use this capability to generate the huge amounts of test data needed from the perspectives of the 2 cameras, which point to a specific object in the scene. This enables the authors to use and train a neural network, since shooting the amount of pictures needed by hand would take too long to consider this idea.

3.2.1 Blender

Besides 3D-modelling Blender enables the user to perform various different actions, such as laying out scenes, UV-Editing, shading, animating and rendering. Additionally scenes can

be modified by executing Python scripts. Figure 3.1 shows the interface of Blender 2.8 with the default file loaded.




Figure 3.1: Blender interface after startup. The default scene features a grey cube, a camera (left of the cube) and a light source (black dot and circle positioned on the top right of the cube).

The authors decided to extensively use the scripting function in their work. One example of a Python script, that can be executed in blender is the following code:

```

1 import bpy
2
3 # Selects all cubes and deletes them
4 bpy.ops.object.select_all(action='DESELECT')
5 bpy.ops.object.select_by_type(type='MESH')
6 bpy.ops.object.delete()
7
8 # Adds a new cube
9 bpy.ops.mesh.primitive_cube_add(size=3, enter_editmode=False, location=(4, 2, 0))
10
11 # Adds a new material representing the colour red
12 bpy.ops.material.new()
13 material = bpy.data.materials[-1]
14 material.name = 'Red'
15 material.diffuse_color = (0.8, 0.1, 0.1, 1)
16

```

```
17 # Apply material onto the newly created cube object  
18 bpy.context.active_object.data.materials.append(material)
```

This code first clears the scene from other meshes (because running the script twice would place the new cube inside the old cube). Then it adds a new mesh in form of a cube at the given location. Next we want to add some color to the cube. For this to work a material is needed, which is basically a specification of how the surface of the object will look like. Advanced materials can represent raw or reflective surfaces, but in order to keep it simple this material will just represent a red surface (represented in red/green/blue/alpha channels ranging from 0.0 (for 0) to 1.0 (for 255)). Lastly the created material is applied to the object. The final result can be seen below:




Figure 3.2: Result of the above code. The default cube is deleted, a new one is added with a red material and translated via the 'location' argument.).

Using similar Python scripts the authors will generate the image data necessary as well as perform the calculation of the distance which will be specified as the correct values to train the neural network on.

3.3 Image preprocessing

After the test images have been rendered with the help of Blender some image preprocessing is required. For example the machine learning component of this project should take two

images as an input. To simplify the input the two images will simply be placed next to each other to form a new image twice as wide as the original images. This process is visualized in Figure 3.3. Other preprocessing measures that have to be taken are downscaling the image, as to not have too many weights in the first layer of the neural network.




Figure 3.3: Two greyscale renders showing the same object from two different camera positions in Blender are merged into one picture which can be feed into the neural network component.

Additionally, the authors decided to experiment if manipulating these test images further would result in differences of distance perception by the neural network or if they would speed up or slow down the learning phase. These manipulations are done using OpenCV and Python3. OpenCV provides many image manipulation tools. For this project the authors chose the following methods of image manipulation:

Name	Description
Greyscale	A greyscale image is an image with only one value for the red, green and blue colour channels, resulting in different shades of grey instead of usual colours.
Resolution	Resolution refers to the number of pixels placed in each dimension (width and height).
Cropping	When cropping an image an unwanted part located at the peripheral areas of the image is removed.
Saturated	Saturated images feature stronger colours, which makes them easier to distinguish from another.
Brightness	Brightening images can make colours harder to distinguish from another. Additionally it can lead to the same problems encountered in overexposed images, such as part of the image being completely white and therefore not providing any information.

3.4 Neural Network

A Neural Network consists of nodes, each receiving some input values as well as some weights associated to each input value and outputs some output value. The calculation returning the output value is relatively simple. Many of these nodes form what is called a layer. A vanilla Neural Network consists of multiple layers, where each node receives all output values of all nodes in the previous layer as an input. A visualization of a Neural Network can be seen in Figure 3.4. By manipulating the weights associated to each input value the network can learn to solve some given task.




Figure 3.4: Depiction af a vanilla Neural Network. Each circle represents a node and the arrows show the flow of information. All nodes on the same level are collectively called a layer.

3.4.1 Convolutional Neural Network

One of the disadvantages of Neural Networks is that they need huge amounts of data in order to optimize the values of all weights. Therefore a modified version of Neural Networks was invented. Convolutional Neural Networks simplify the optimization of the weights by stating that some weights are shared between multiple connections between nodes. This results in fewer weights having to be optimized. Additionally Convolutional Neural Networks contribute to the fact that often problems require to perform similar actions in multiple parts of the input data, e.g. search for edges in all sections of an image.

3.4.2 Loss Function

A loss function is used in optimization problems to determine the loss or cost of some operation. Therefore a low loss is desired. Some loss functions commonly used loss functions are:




Figure 3.5: Segment of a Convolutional Neural Network. In comparison to a vanilla Neural Network each node is only connected to its closest neighbours. Additionally all arrows connected to another node in the same relative position (same colour) have the same weight attached.

Mean Squared Error

$$mse = \frac{1}{n} \sum_{i=1}^n (x_i - \hat{x}_i)^2$$

where n is the number of elements, x is the correct value and \hat{x} is the estimated value.

Mean Absolute Error

$$mae = \frac{\sum_{i=1}^n |\hat{x}_i - x_i|}{n}$$

where n is the number of elements, x is the correct value and \hat{x} is the estimated value.

Mean Absolute Percentage Error

$$mape = \frac{100\%}{n} \sum_{i=1}^n \left| \frac{x_i - \hat{x}_i}{x_i} \right|$$

where n is the number of elements, x is the correct value and \hat{x} is the estimated value.

Mean Squared Logarithmic Error

$$msle = \frac{\sum_{i=1}^n \log(\hat{x}_i + 1) - \log(x_i + 1)}{n}$$

where n is the number of elements, x is the correct value and \hat{x} is the estimated value.

Logcosh

Kullback Leibler Divergence

3.4.3 Activation Function

The activation function defines the output of a node based on the input values it receives. Some of the most widely used activation functions are displayed from Figure 3.6 to Figure 3.14.




Figure 3.6: ...

3.5 Neural Network Implementations

The authors decided to use a software framework called TensorFlow for the first implementation of the neural network. This has the following two advantages: using Tensorflow allows for a low effort proof of concept and it makes testing out different configurations (e.g. number of hidden layers or filters in image preprocessing) of the neural network easier.

After it has been shown that the challenge of detecting the distance to an object can be solved using machine learning, the authors plan on implementing a neural network in C++ on their own. The knowledge gained in the TensorFlow implementation will be used in the C++ implementation, which hopefully will make the work less time consuming.




Figure 3.7: ...




Figure 3.8: ...




Figure 3.9: ...




Figure 3.10: ...

3.6 TensorFlow

As machine learning has gained popularity in recent years the demand for applicable frameworks grew. One of the most popular is called TensorFlow. It was developed by Google for




Figure 3.11: ...




Figure 3.12: ...

internal use and was published under the Apache License 2.0 on the 9th of November 2015. TensorFlow supports APIs for Python, C, C++, Go, Java, JavaScript and Swift. Due to its




Figure 3.13: ...




Figure 3.14: ...

popularity third party APIs for C#, R, Scala, Rust and many more were developed.
Its use cases reach from categorizing handwritten digits to YouTube video recommendations,

one of the many applications Google use it for.

Tom Hope et al. describe TensorFlow as a software framework for numerical computations based on dataflow graphs¹.

3.6.1 Computation graph

To compute a value using TensorFlow a computation graph has to be constructed. In this graph each node corresponds to an operation, such as subtraction or division. By connecting these nodes via edges the output of one node can be fed as input into another node. One example of such a computation graph can be seen in Figure 3.15.




Figure 3.15: Each node represents an operation, where *const* stands for a constant value, *add* for addition, *mul* for multiplication and *div* for division. Edges, represented by arrows, connect nodes. The information shared between the nodes is described by the numbers written next to the edges. This computational graph calculates the result of the arithmetic expression $(28 * 2) + 28/2$.

The implementation of the computation graph, shown in Figure 3.15, in Python could look as follows:

```
1 import tensorflow as tf
2
3 a = tf.constant(28)
4 b = tf.constant(2)
5
6 c = tf.multiply(a, b)
7 d = tf.add(a, c)
8 e = tf.divide(d, b)
9
10 with tf.Session() as sess:
11     out = sess.run(e)
12
13 print(out)
```

¹Itay Lieder Tom Hope Yehezkel S. Resheff. *Learning TensorFlow*. O'Reilly, 2017, page 6.

The first line specifies that the TensorFlow functionality should be imported. Line 3 and 4 define the two constant values and assigns them the values 28 and 2 respectively. In line 6 to 8 the other nodes of the graph are specified. E.g. in line 6 a new node, named c , is created and the output of node a and node b are connected as its input. To perform the calculation described by the graph a new session is created in line 10. Finally the output of the graph (node e) is specified in line 11, the result is calculated and printed in line 13.

TensorFlow allows for another way of specifying a graph with these arithmetic operations:

```
1 import tensorflow as tf
2
3 a = tf.constant(28)
4 b = tf.constant(2)
5
6 e = (a * b + a) / b
7
8 with tf.Session() as sess:
9     out = sess.run(e)
10
11 print(out)
```

This code is equivalent to the first one, but uses syntactic sugar to shorten line 6 to 8 in the first code block into line 6. At this point it should be noted that while it might look like it line 6 does not calculate anything. It simply describes how the computational graph should look. The answer (42) is calculated in the session in line 9.

Using the same principle, but different functions TensorFlow allows the creation of complex Neural Networks in a simple fashion.

3.6.2 Alternatives to Tensorflow

TensorFlow offers some abstraction libraries, such as Keras. As some code snippets of are used very frequently and with only slight variation when implementing a Neural Network, Keras offers these code blocks as predefined functions, making development easier for “standard” use cases.

For example Keras provides several different layers, that can be used to put together a Neural Network. Table 3.6.2 describes some of the available layers.

[INFO: Abstraction libraries such as Keras and TF-Slim offer simplified high-level access to the ”LEGO bricks” in the lower-level library, helping to streamline the construction of the dataflow graphs, training them, and running inference.^{4]}

[+ warum verwenden wir ausgerechnet Tenserflow]

⁴Tom Hope, *Learning TensorFlow*, page 7.

Name	Description
Dense	regular densely connected layer
Activation	Applies an activation function to an output
Dropout	Applies Dropout to the input
Flatten	Flattens the input down to shape (None, <number of elements>)
Conv1D	1D convolution layer (e.g. temporal convolution)
MaxPooling1D	Max pooling operation for temporal data
LocallyConnected1D	Locally-connected layer for 1D inputs
SimpleRNN	Fully-connected RNN where the output is to be fed back to input

Table 3.1: Short description of layers available in Keras, that can be used to assemble a Neural Network, with little effort³.

	Open Source	Actively developed	Parallelization	Interface
TensorFlow	Yes	Yes	Yes	Python, C, C++, Go, Java, JavaScript, Swift, R, Julia
Keras	Yes	Yes	Yes	Python
PyTorch	Yes	Yes	Yes	Python, C++
Torch	Yes	No	Yes	Lua, LuaJIT, C, C++, OpenGL
Wolfram Mathematica	No	Yes	Yes	Wolfram Language

Chapter 4

ROS2

Author:

4.1 What is ROS?

[Ubuntu?]

[INFO: ROS (Robot Operating System) provides libraries and tools to help software developers create robot applications.¹]

[TEST: asdf²] [Core Concepts]

4.1.1 Nodes

[Wie verwenden wir das? / Wie macht es unsere Arbeit leicher?]

4.1.2 ...

4.2 Why use ROS?

4.3 Comparing ROS and ROS2

[python2 < python3]

¹Inc. Open Source Robotics Foundation. *Documentation - ros wiki*. 2019. URL: <https://wiki.ros.org>.

²Open Source Robotics Foundation, *Documentation - ros wiki*.

Chapter 5

Implementation

Author:

5.1 Generating test data

Good test data is of utmost importance in machine learning. The system can only know information that is depicted in the training data, which is why it is important to include as many aspects of the problem as possible in this data.

Since machine learning needs a lot of data in order to solve the given task it can be tiresome to generate and label all this data by hand. Therefore the authors decided to simulate the objects and the camera using a computer graphics modelling software called Blender.

Blender allows for relatively easy generation of training data by providing a Python API. With this API almost anything that can be done using the blender user interface can also be done using Python.

For generating data the first step for the authors was to model a scene, as seen in Figure 5.1. This scene would contain a building containing rooms. Each of the rooms is home to one object. And this object was rendered using two cameras, representing two points of view. The first camera is positioned relative to the object, whereas the second camera is positioned relative to the first camera, guaranteeing a somewhat different view on the object. Lastly, the two renders formed one input for the neural network.

For the test objects the authors used two objects. One is a already existing monkey head in blender, the other one has been self-modelled and represents a vase with some decorating sticks in it. Both objects can be seen in Figure 5.2 and in Figure 5.3.

5.2 OpenCV

OpenCV is a framework for image manipulation. Some of its use cases are changing the colour spectrum, filtering the image by colour and cropping images. The authors use OpenCV to test whether there are differences between filters for the images in the training data, for example greyscale images compared to coloured images. An example render, which OpenCV gets as an input can be seen in Figure 5.4.




Figure 5.1: The test scene for generating data. Each object is placed in a room and is separated by walls.




Figure 5.2: One of the test objects - a monkey head.

5.2.1 Greyscale

Converting an image into greyscale can easily be achieved by the following OpenCV code:
[TODO: How is the image converted? (Average of all three channels?, ...)]

```
1 import cv2
2
3 image = cv2.imread('path/to/image')
```

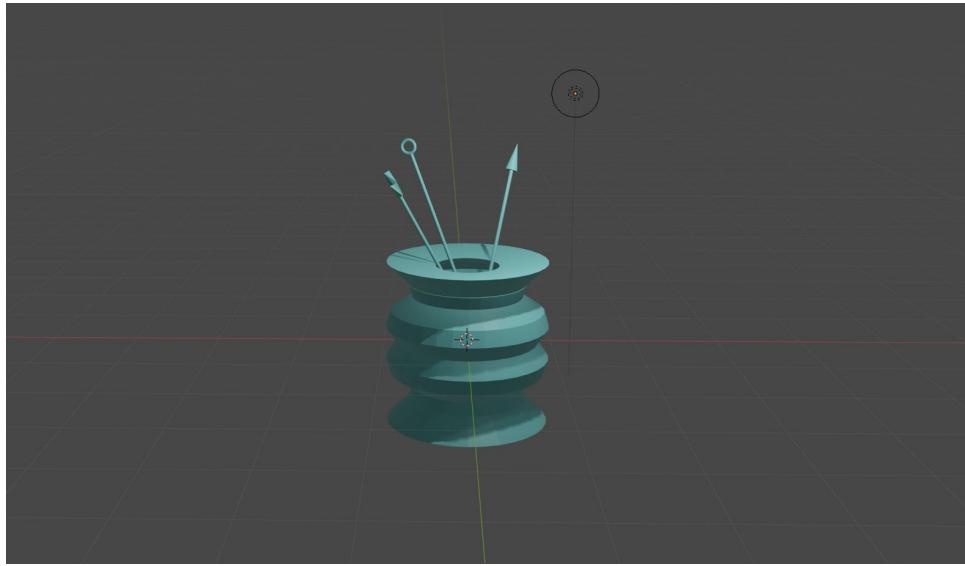


Figure 5.3: One of the test objects - a decorative vase.



Figure 5.4: One of two original renders produced by Blender, both depicting the same object from different points of view.

```
4 image_greyscale = cv2.cvtColor(image, cv2.COLOR_BGR2GRAY)
5 cv2.imwrite('path/for/saving/greyscale/image', image_greyscale)
```

This code first reads the image into 'image'. Then it converts the color of 'image' into a greyscale format and stores the result into 'image_greyscale', which is then written to the specified path. The output generated by this code is depicted in Figure 5.5.




Figure 5.5: The greyscale image produced by OpenCV.

The advantage of using greyscale images in Neural Networks is the simplification of the input layer. In greyscale images each pixel can be represented by a single decimal value between 0 and 1. This enables the first layer of the neural network to be two dimensional instead of the three dimensional counterpart, where each pixel is represented by the three decimal values for the red, green and blue colour channels.

5.2.2 Resolution

By reducing the resolution of an image the density of pixels is lessened. In this process information, that can not be regained, is lost. The OpenCV code for downscaling the images used by the authors is the following:

```
1 import cv2
2
3 image = cv2.imread('path/to/image')
4
5 scale_percent = 10 # percent of original size
6 width = int(image.shape[1] * scale_percent / 100)
7 height = int(image.shape[0] * scale_percent / 100)
8 dim = (width, height)
9
10 downscaled = cv2.resize(image, dim)
11 cv2.imwrite('{}{}'.format(newdir_path, filename), downscaled)
```

OpenCV provides a `resize` function, which takes an image and the new dimensions of the image as an argument and outputs the resized image. To make sure the aspect ratio stays the same new image dimensions are calculated as a percentage of the original ones.

The output image of this code can be found in Figure 5.6.




Figure 5.6: Image with lower resolution than the original one. Due to the fact that the image is displayed in the same size as Figure 5.4, the pixel size in this image appears larger.

Downscaling the images before they are passed into the Neural Network can be profitable, because the number of weights in the first layer of the network is reduced. This can advance the learning speed.

5.2.3 Cropping

Cropping an image can remove unnecessary or unwanted parts of an image by simply cutting off areas. This is often wanted in photography to only keep what is interesting in a photo and to shift the view of the viewer to specific areas. However, the authors guess that cropping an image will worsen performance of the neural network, because information of the relative sizes are partly lost. It basically compares to zooming into the image.

Such a cropped image produced by the following code can be found in Figure 5.7.

```
1 import cv2
2
3 image = cv2.imread('path/to/image')
4
5 crop_margin_percent = 5
6 crop_margin_width = int(image.shape[1] * crop_margin_percent / 100)
7 crop_margin_height = int(image.shape[0] * crop_margin_percent / 100)
8
9 crop_img = image[crop_margin_width:-crop_margin_width, crop_margin_height:-
    crop_margin_height]
10 cv2.imwrite('{}/{}'.format(newdir_path, filename), crop_img)
```



Figure 5.7: In this image five percent of each side was removed, therefore the object appears closer if the image is displayed in the same size. Some information positioned in the outer areas of the image are lost during the cropping process.

5.2.4 Saturated

A saturated image means stronger colours, basically making them more distinguishable from each other. If an image is not saturated enough, colours appear as "washed out" and differences in colour are difficult to determine. Therefore, in order to help the neural network, the authors decided to also test with saturated versions of these images.

```
1 import cv2
2
3 image = cv2.imread('path/to/image')
4
5 hsv = cv2.cvtColor(image, cv2.COLOR_BGR2HSV).astype('float32')
6 (h, s, v) = cv2.split(hsv)
7 s *= 1.5
8 s = np.clip(s, 0, 255)
9 hsv = cv2.merge([h, s, v])
10 saturated = cv2.cvtColor(hsv.astype('uint8'), cv2.COLOR_HSV2BGR)
11 cv2.imwrite('{}{}'.format(newdir_path, filename), saturated)
```

To saturate an image using OpenCV it has to be converted into the HSV colour representation. The HSV representation specifies each colour as a hue, a saturation and a value. Therefore changing the saturation in this model is relatively easy, as the saturation value of each pixel can simply be multiplied by a constant. Before the image can be converted back into the bgr colour representation the saturation values are restricted between 0 and 255, the lowest and highest saturation possible. The result achieved by this code can be seen in Figure 5.8.



Figure 5.8: Comparison between normal and saturated image.

5.2.5 Brightness

Because the neural network should also work in different environment, where other methods could have problems, the authors also tested with overly bright images. This makes it very hard to notice dark areas, such as shadows, to help with distances between objects and scaling of objects.

```
1 import cv2
2
3 image = cv2.imread('path/to/image')
4
5 alpha = 1 # contrast
6 beta = 60 # brightness
7 bright_img = cv2.convertScaleAbs(image, alpha=alpha, beta=beta)
8
9 cv2.imwrite('{}/{}'.format(newdir_path, filename), bright_img)
```

5.3 Neural Network

5.3.1 Structure of our Neural Network

[wie lesen wir Daten ein, wie viele layer, was ist der output (maximale entfernug? z.B. 10m)]



Figure 5.9: The original image (Figure 5.4) modified by the brightening code. The image appears more washed out, as all colours are move similar due to them having moved closer to white.

5.4 C++ Implementation

5.5 Technical difficulties

Chapter 6

Experiment 1

Author:

6.1 Environment

[TODO: viele Fotos] [TODO: Hintergrundfarbe, Untergrundfarbe, Struktur (Hintergrund und Untergrund), Beleuchtung (Art, Helligkeit, Richtung, mehrere Lichtquellen, welche Lichtquellen, ...)]

6.2 Setup

[TODO: viele Fotos] [TODO: welche Objekte (Größe, Farbe, wie viele (mindestens 3)?, ...), Kameras, Entfernung zu Objekten]

6.3 Sequence of Events

6.4 Results

Chapter 7

Lessons learned

Author:

Chapter 8

Experiment 2

Author:

8.1 Environment

[TODO: viele Fotos]

8.2 Setup

[TODO: viele Fotos]

8.3 Materials

[TODO: viele Fotos]

8.4 Sequence of Events

8.5 Results

Chapter 9

Conclusion

Author:

Index

Authors Index

Bresenham, Jack, 7

K.T. Gribbon C.T. Johnston, Donald G.
Bailey, 6

Open Source Robotics Foundation, Inc.,
23

SAS, Parrot Drones, 3

Tom Hope Yehezkel S. Resheff, Itay
Lieder, 20, 21

Zezhi Chen Chengke Wu, Hung Tat Tsui,
7

Literature Index

A linear algorithm for incremental digital display of circular arcs, 7

A new image rectification algorithm, 7

A real-time FPGA implementation of a barrel distortion correction algorithm with bilinear interpolation, 6

Documentation - ros wiki, 23

Learning TensorFlow, 20, 21


Parrot Bebop 2 Specifications, 3

Bibliography

- Bresenham, Jack. “A linear algorithm for incremental digital display of circular arcs”. In: *Communications of the ACM* 20.2 (1977), pp. 100–106.
- K.T. Gribbon C.T. Johnston, Donald G. Bailey. “A real-time FPGA implementation of a barrel distortion correction algorithm with bilinear interpolation”. In: *Image and Vision Computing New Zealand*. 2003, pp. 408–413.
- Keras Documentation*. 2019. URL: <https://keras.io/>.
- Open Source Robotics Foundation, Inc. *Documentation - ros wiki*. 2019. URL: <https://wiki.ros.org>.
- SAS, Parrot Drones. *Parrot Bebop 2 Specifications*. 2019. URL: <https://www.parrot.com/us/drones/parrot-bebop-2>.
- Tom Hope Yehezkel S. Resheff, Itay Lieder. *Learning TensorFlow*. O'Reilly, 2017.
- Zezhi Chen Chengke Wu, Hung Tat Tsui. “A new image rectification algorithm”. In: *Pattern Recognition Letters* 24.1-3 (2003), pp. 251–260.

Messbox zur Druckkontrolle

— Druckgröße kontrollieren! —



Breite = 100 mm
Höhe = 50 mm

— Diese Seite nach dem Druck entfernen! —